

# Introduction to Big Data Analysis Ensemble Methods

Zhen Zhang

Southern University of Science and Technology

# Outlines

Introduction

Bagging and Random Forest

Boosting and AdaBoost

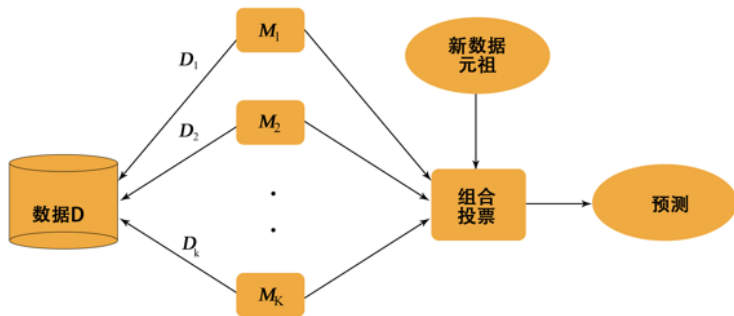
Gradient Boosting Decision Tree

XGBoost

Conclusion and Python Examples

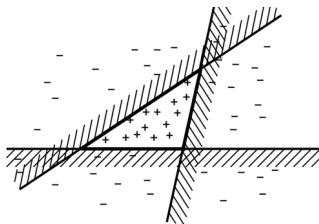
## Ensemble Methods

- Wisdom of Crowds (“三个臭皮匠，顶个诸葛亮”)
- Multiple weak learners (base learners, may be heterogenous) can improve learning performance



## Why it can improve the performance

- More flexible to interpret
  - Single linear classifier (perceptron) does not work
  - Try multiple classifiers



- Reduce misclassification rate
  - Misclassification rate of single classifier is  $p$
  - Choose  $N$  classifiers, same but independent, voting
  - Error rate of majority vote =  $\sum_{k > N/2} \binom{N}{k} p^k (1-p)^{N-k}$
  - When  $N = 5, p = 0.1$ , Error rate  $< 0.01$


## Two commonly used ensemble methods

- Bagging

- Random sampling : generating independent models, and averaging for regressions (making majority vote for classifications)
- Reducing variances
- Example : Random forests

- Boosting

- Sequential training : training the subsequent models based on the errors of previous models
- Reducing bias
- Examples : AdaBoost and GBDT



# Outlines

Introduction

Bagging and Random Forest

Boosting and AdaBoost

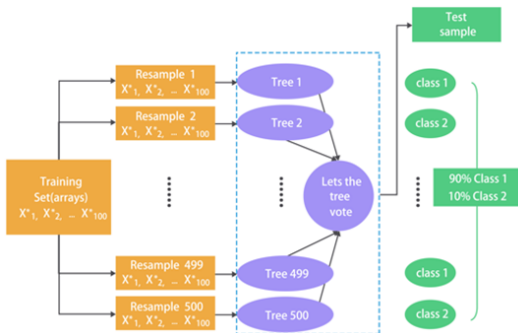
Gradient Boosting Decision Tree

XGBoost

Conclusion and Python Examples

# Bagging

- Bagging is short for bootstrap aggregation
- Bagging generates a committee of predictors and combine them in a certain manner to the final model
- Single predictor suffers from instability, while bagging could improve the stability by majority vote (classification) or averaging (regression) over all single predictors



# Sampling

- Given a dataset  $D$  of  $n$  samples, at the iteration  $m = 1, \dots, M$ , the training set  $D_m$  is obtained by sampling from  $D$  with replacement. Then  $D_m$  is used to construct classifier  $\hat{f}_m(x)$ .
- Sampling with replacement : some samples in  $D$  may be missing in  $D_m$ , while some other samples may occur more than once
- On average, 63.2% of the samples in  $D$  could be selected into  $D_m$ . In fact, for each sample, the probability that it is not selected in one round is  $1 - \frac{1}{n}$ . Then it is not selected in all  $n$  rounds with probability  $\lim_{n \rightarrow \infty} (1 - \frac{1}{n})^n = 0.368$ .



# Algorithm

- Input : training set  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$
  - Output : additive model  $\hat{f}_{bag}(x)$
1. For  $m = 1$  to  $M$  :
    - 1.1 Sample from  $D$  with replacement to obtain  $D_m$
    - 1.2 Train a model  $\hat{f}_m(x)$  from the dataset  $D_m$  : for classification,  $\hat{f}_m(x)$  returns a K-class 0-1 vector  $e_k$  ; for regression, it is just a value
  2. Compute bagging estimate  $\hat{f}_{bag}(x) = \frac{1}{M} \sum_{m=1}^M \hat{f}_m(x)$  : for classification, make majority vote  $\hat{G}_{bag}(x) = \arg \max_k \hat{f}_k(x)$  ; for regression, just return the average value

## Variance Reduction

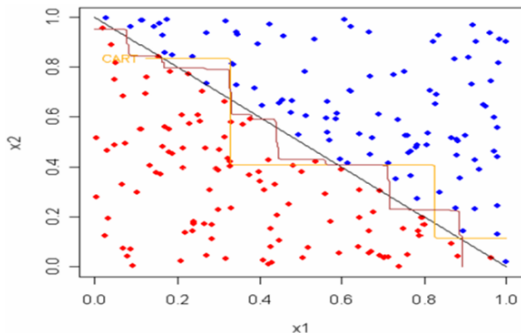
- In bagging, we use the same model to train different sample set in each iteration ; assume the models  $\{\hat{f}_m(x)\}_{m=1}^M$  have the same variance  $\sigma^2(x)$ , while the correlation of each pair is  $\rho(x)$
- Then the variance of the final model is :

$$\begin{aligned}\text{Var}(\hat{f}_{bag}(x)) &= \frac{1}{M^2} \left( \sum_{m=1}^M \text{Var}(\hat{f}_m(x)) + \sum_{t \neq m} \text{Cov}(\hat{f}_t(x), \hat{f}_m(x)) \right) \\ &= \rho(x)\sigma^2(x) + \frac{1 - \rho(x)}{M}\sigma^2(x)\end{aligned}$$

- As  $M \rightarrow \infty$ ,  $\text{Var}(\hat{f}_{bag}(x)) \rightarrow \rho(x)\sigma^2(x)$ . This usually reduces the variance.
- If  $\rho(x) = 0$ , the variance could approach zero
- The random sampling in bagging is to reduce the correlation  $\rho(x)$ , i.e., make the sub-predictors as independent as possible







# Limitations of Decision Tree

- Stuck at local optimum : The greedy algorithm makes it stop at the local optimum, as it seeks the maximal information gain in each tree split
- Decision boundary : Use one feature in each split, the decision boundary is parallel to the coordinate axes
- Bad interpretability and instability




## Bagging Tree vs. Single Tree

- The true split point is  $x_1 = 0.5$  and the Bayes error is 0.2 ( $P(Y = 1|x_1 \leq 0.5) = 0.2$  and  $P(Y = 1|x_1 > 0.5) = 0.8$ ). # of features is  $p=5$  and # of samples is  $n=2000$ .



## Error of Bagging Trees

- Bagging trees reduce the test error by reducing variance, leaving bias unchanged
- Consensus : majority vote ; Probability : average the probability predicted by sub-models



# Random Forest

- Random Forest further reduces the variance by adding independency to the committee of decision trees
- This is achieved by introducing more randomness.
- Three level of randomness :
  - Sampling on the training data with replacement
  - Select features at random
  - Select split point at random
- No pruning is needed.
- Example : RF consisting of 3 independent trees, each with an error rate of 40%. Then the probability that more than one tree misclassify the samples is
$$0.4^3 + 3 * 0.4^2 * (1 - 0.4) = 0.352$$

# Random Forest Algorithm

- Input : training set  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$
- Output : additive model  $\hat{f}_{rf}(x)$

1. For  $m = 1$  to  $M$  :

1.1 Sample from  $D$  with replacement to obtain  $D_m$

1.2 Grow a random-forest tree  $T_m$  to the dataset  $D_m$  : by recursively repeating the following steps for each terminal node of the tree, until the minimum node size  $n_{min}$  is reached

1.2.1 Select  $q$  features at random from the  $p$  features

1.2.2 Pick the best feature/split-point among the  $q$

1.2.3 Split the node into two daughter nodes

2. Output the ensemble of trees  $\{T_m\}_{m=1}^M$  : for regression,

$$\hat{f}_{rf}(x) = \frac{1}{M} \sum_{m=1}^M T_m(x) : \text{for classification, make majority vote}$$

- Small value of  $q$  increases the independency of trees ; empirically,  $q = \log_2 p + 1$

## Model Evaluation

- Margins : The difference between the percentage of decision trees that correctly classify the samples and the percentage of trees misclassifying the samples
- Out-of-bag (OOB) erros : The observation is called out-of-bag sample to some trees if it is not sampled for those trees. Denote the training set in the  $m$ -th sampling by  $D_m$ . OOB error is computed as :

1. For each observation  $(x_i, y_i)$ , find the trees which treat it as OOB sample :  $\{\hat{T}_m(\mathbf{x}) : (\mathbf{x}_i, y_i) \notin D_m\}$
2. Use those trees to classify this observation and make majority vote as the label of this observation :


$$\hat{f}_{oob}(\mathbf{x}_i) = \arg \max_{y \in \mathcal{Y}} \sum_{m=1}^M \mathbb{I}(\hat{f}_m(\mathbf{x}_i) = y) \mathbb{I}(\mathbf{x}_i \notin D_m)$$

3. Compute the number of misclassified samples, and take the ratio of this number to the total number of samples as OOB error :  $Err_{oob} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{f}_{oob}(\mathbf{x}_i) \neq y_i)$



# Feature Importance

- Using split criteria
  - The improvement in the split-criterion as feature importance
  - It is accumulated over all the trees for each variable
- Using OOB randomization
  - Randomly permute the values of each feature in the OOB samples, and compute the prediction accuracy
  - The decrease in accuracy as a result of this permutation is averaged over all trees as feature importance



## Proximity Matrix

- For every tree, any pair of OOB observations sharing a terminal node has their proximity increased by one
- The more two OOB samples sharing the same node, the more similar the two samples are
- Proximity matrix could be used to define distance between two samples

# Pros and Cons

- Where it is good
  - Bagging or random forest (RF) work for models with high variance but low bias
  - Better for nonlinear estimators
  - RF works for very high-dimensional data, and no need to do feature selection as RF gives the feature importance
  - Easy to do parallel computing
- Disadvantage
  - Overfitting when the samples are large-sized with great noise, or when the dimension of data is low
  - Slow computing performance comparing to single tree
  - Hard to interpret

# Outlines

Introduction

Bagging and Random Forest

**Boosting and AdaBoost**


Gradient Boosting Decision Tree

XGBoost

Conclusion and Python Examples

# Boosting

- Boosting : combines the outputs of many “weak” classifiers to produce a powerful “committee”
- Weak classifier : error rate  $< 0.5$  (random guessing)
- **Sequentially** apply the weak classifiers to the repeatedly modified data, emphasizing the misclassified samples
- Combine weak classifiers through a weighted majority vote or averaging to produce the final prediction



## Boosting Fits an Additive Model

- Additive model :  $f(x) = \sum_{m=1}^M \beta_m b(x; \gamma_m)$
- Possible choices for basis function  $b(x; \gamma)$  :
  - Neural networks :  $\sigma(\gamma_0 + \gamma_1^T x)$ , where  $\sigma(t) = 1/(1 + e^{-t})$
  - Wavelets
  - Cubic spline basis
  - Trees
  - Eigenfunctions in reproducing kernel Hilbert space (RKHS)
- Parameter fitting :  $\min_{\{\beta_m, \gamma_m\}} \sum_{i=1}^N L(y_i, \sum_{m=1}^M \beta_m b(x_i; \gamma_m))$
- Loss function : squared error  $L(y, f(x)) = (y - f(x))^2$  or likelihood-based loss

## Forward Stagewise Additive Modeling

- Input : training set  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$
- Output : additive model  $f_M(x)$

1. Initialize  $f_0(x) = 0$

2. For  $m = 1$  to  $M$  :

2.1 Compute  $(\beta_m, \gamma_m) = \arg \min_{\beta, \gamma} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + \beta b(x_i; \gamma))$

2.2 Update  $f_m(x) = f_{m-1}(x) + \beta_m b(x; \gamma_m)$

- Squared error loss : in step 2.1,

$$L(y_i, f_{m-1}(x_i) + \beta b(x_i; \gamma)) = \underbrace{(y_i - f_{m-1}(x_i))}_{\text{residual}} - \beta b(x_i; \gamma)^2$$

## Exponential Loss and AdaBoost

- Exponential loss :  $L(y, f(x)) = \exp(-yf(x))$
- Classifier as basis function :  $b(x; \gamma) = G(x) \in \{-1, 1\}$
- Let  $w_i^{(m)} = \exp(-y_i f_{m-1}(x_i))$ , then step 2.1 turns to be :

$$\begin{aligned}
 (\beta_m, G_m) &= \arg \min_{\beta, G} \sum_{i=1}^n w_i^{(m)} \exp(-\beta y_i G(x_i)) \\
 &= \arg \min_{\beta, G} \left[ \sum_{y_i \neq G(x_i)} w_i^{(m)} (e^\beta - e^{-\beta}) + e^{-\beta} \sum_{i=1}^n w_i^{(m)} \right]
 \end{aligned}$$

- $G_m = \arg \min_G \sum_{i=1}^n w_i^{(m)} I(y_i \neq G(x_i)).$
- $\beta_m = \arg \min_{\beta} \left[ \epsilon_m (e^\beta - e^{-\beta}) + e^{-\beta} \right] = \frac{1}{2} \log \frac{1-\epsilon_m}{\epsilon_m}$  where  
 $\epsilon_m = (\sum_{i=1}^n w_i^{(m)} I(y_i \neq G(x_i))) / \sum_{i=1}^n w_i^{(m)}$  is weighted error rate




# AdaBoost Algorithm

- Input : training set  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$ , loss function  $L(y, f(x))$
  - Output : Weighted classifier  $G(x)$
1. Initialize  $w_i = 1/N$ ,  $i = 1, \dots, N$
  2. For  $m = 1$  to  $M$  :
    - 2.1 Fit a classifier  $G_m(x)$  to the training data  $D$  with weight  $\{w_i\}$
    - 2.2 Compute the error  $\epsilon_m = (\sum_{i=1}^n w_i^{(m)} I(y_i \neq G(x_i))) / \sum_{i=1}^n w_i^{(m)}$
    - 2.3 Compute  $\alpha_m = \log \frac{1-\epsilon_m}{\epsilon_m}$  ( $\alpha_m = 2\beta_m > 1$ )
    - 2.4 Update the weight  $w_i^{(m+1)} = w_i^{(m)} \exp(\alpha_m I(y_i \neq G_m(x_i)))$ , for  $i = 1, \dots, N$
  3. Output  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$


# Illustration

- Weights of weak classifiers :  
the better the classifier is,  
the larger its weight is
- Weights of samples :  
Re-weighting after each  
step, increase the weights  
for misclassified samples
- Simulation : 2-class  
classification, 1000 training  
samples from each class,  
10,000 test samples ;  
two-leaf classification tree  
(stump) as base learner



# Loss Functions

- For classification, exponential loss and binomial negative log-likelihood (deviance) loss  $\log(1 + \exp(-2yf))$  share the same population minimizer; thus it is equivalent to MLE rule
- For classification, squared error loss is not good (not monotonically decreasing); the exponential loss is good and binomial deviance is better (less penalty for large  $-yf$ )



# Pros and Cons

- Where it is good
  - AdaBoost improve the classification performance comparing to weak classifiers
  - Many choices for weak classifiers : trees, SVMs, kNNs, etc.
  - Only one tuning parameter  $M$  : # of weak classifiers
  - prevent overfitting suffered by single weak classifiers (e.g. complex decision tree)
- Disadvantage
  - Weak interpretability
  - Overfitting when using very bad weak classifiers
  - Sensitive to outliers
  - Not easy for parallel computing

# Outlines

Introduction

Bagging and Random Forest

Boosting and AdaBoost

**Gradient Boosting Decision Tree**

XGBoost

Conclusion and Python Examples

## Boosting Tree

- Using classification trees or regression trees as base learners
- $f_M(x) = \sum_{m=1}^M T(x; \Theta_m)$  where  $T(x; \Theta) = \sum_{j=1}^J \gamma_j I(x \in R_j)$
- Parameter set  $\Theta = \{R_j, \gamma_j\}_{j=1}^J$
- Parameter finding : minimizing the empirical risk

$$\hat{\Theta} = \arg \min_{\Theta} \sum_{j=1}^J \sum_{x_i \in R_j} L(y_i, \gamma_j) \quad (\text{Combinatorial optimization})$$

- Approximate suboptimal solutions :
  1. Finding  $\gamma_j$  given  $R_j$  :  $\gamma_j = \bar{y}_j = \frac{1}{|R_j|} \sum_{y_i \in R_j} y_i$  for  $L^2$  loss; and  
 $\gamma_j = \text{modal class in } R_j$  for misclassification loss
  2. Finding  $R_j$  given  $\gamma_j$  : Difficult, need to estimate  $\gamma_j$  as well ;  
 greedy, top-down recursive partitioning algorithm

## Boosting Tree as Forward Stagewise Algorithm

- $\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + T(x_i; \Theta_m))$ 
  1.  $\hat{\gamma}_{jm} = \arg \min_{\gamma_{jm}} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma_{jm})$
  2. Finding  $R_{jm}$  is more difficult than for a single tree in general.
- Squared-error loss : fit a tree to the residual
 
$$L(y_i, f_{m-1}(x_i) + T(x_i; \Theta_m)) = \underbrace{(y_i - f_{m-1}(x_i))}_{\text{residual}} - T(x_i; \Theta_m))^2$$
- Two-class classification and exponential loss : AdaBoost for trees,  $\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^N w_i^{(m)} \exp[-y_i T(x_i; \Theta_m)]$ 
  1.  $\hat{\gamma}_{jm} = \log \frac{\sum_{x_i \in R_{jm}} w_i^{(m)} I(y_i=1)}{\sum_{x_i \in R_{jm}} w_i^{(m)} I(y_i=-1)}$
- Absolute error or the Huber loss : robust but slow

# Gradient Descent for General Loss

- Supervised learning is equivalent to the optimization problem

$$\min_{\mathbf{f}} L(\mathbf{f}) = \min_{\mathbf{f}} \sum_{i=1}^N L(y_i, f(x_i))$$

- Numerical optimization :  $\hat{\mathbf{f}} = \arg \min_{\mathbf{f}} L(\mathbf{f})$  where  $\mathbf{f} = \{f(x_1), f(x_2), \dots, f(x_N)\}$ ,
- Approximate  $\hat{\mathbf{f}}$  by  $\mathbf{f}_M = \sum_{m=0}^M \mathbf{h}_m$ , where  $\mathbf{f}_0 = \mathbf{h}_0$  is initial guess
- Gradient descent method :  $\mathbf{f}_m = \mathbf{f}_{m-1} - \rho_m \mathbf{g}_m$ , where  $\mathbf{g}_{im} = \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f(x_i)=f_{m-1}(x_i)}$ , and  $\mathbf{h}_m = -\rho_m \mathbf{g}_m$



# Gradient Boosting Decision Tree (GBDT)

- Find a tree  $T(x; \Theta_m)$  by minimization problem

$$\tilde{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^N (-g_{im} - T(x_i; \Theta_m))^2$$

In general  $\tilde{R}_{jm} \neq R_{jm}$

Setting	Loss Function	$-\partial L(y_i, f(x_i)) / \partial f(x_i)$
Regression	$\frac{1}{2}[y_i - f(x_i)]^2$	$y_i - f(x_i)$
Regression	$ y_i - f(x_i) $	$\text{sign}[y_i - f(x_i)]$
Regression	Huber	$y_i - f(x_i)$ for $ y_i - f(x_i)  \leq \delta_m$ $\delta_m \text{sign}[y_i - f(x_i)]$ for $ y_i - f(x_i)  > \delta_m$ where $\delta_m = \alpha \text{th-quantile}\{ y_i - f(x_i) \}$
Classification	Deviance	$k$ th component: $I(y_i = \mathcal{G}_k) - p_k(x_i)$

## GBDT Algorithm

- Input : training set  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$ , loss function  $L(y, f(x))$
- Output : boosting tree  $\hat{f}(x)$

1. Initialize  $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$

2. For  $m = 1$  to  $M$  :

- 2.1 For  $i = 1, 2, \dots, N$  compute  $r_{im} = \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}$

- 2.2 Fit a regression tree to the target (residual)  $r_{im}$ , giving terminal regions  $R_{jm}, j = 1, \dots, J_m$

- 2.3 For  $j = 1, \dots, J_m$ , compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma)$$


- 2.4 Update  $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x_i \in R_{jm})$

3.  $\hat{f}(x) = f_M(x)$

## Regularization Techniques

- Shrinkage : the step 2.4 is modified as  

$$f_m(x) = f_{m-1}(x) + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x_i \in R_{jm})$$
- Subsampling : at each iteration, sample a fraction  $\eta$  of the training set and grow the next tree using the subsample
- Shrinkage + subsampling : best performance



# Feature importance and Partial Dependence Plots

- Feature importance
  - When fitting a single tree  $T$ , at each node  $t$ , one feature  $X_{v(t)}$  and one separate value  $X_{v(t)} = c_{v(t)}$  are chosen to improve a certain quantity of criterion (e.g. GINI, entropy, squared error, etc.)
  - Sum all these improvements  $i_t$  brought by each feature  $X_k$  over all internal nodes : 
$$I_k(T) = \sum_{t=1}^{J-1} i_t I(v(t) = k)$$
  - Average the improvements of all trees  $\Rightarrow$  importance of that feature : 
$$I_k = \frac{1}{M} \sum_{m=1}^M I_k(T_m)$$
- Partial Dependence Plots
  - Partial dependence of  $f(X)$  on  $X_S$  :  $f_S(X_S) = E_{X_C} f(X_S, X_C)$
  - Estimate by empirical mean : 
$$\bar{f}_S(X_S) = \frac{1}{N} \sum_{i=1}^N f(X_S, x_{iC})$$

# Pros and Cons

- Where it is good
  - For all regression problems
  - Better for two-class classification, possible for multi-class problems (not suggested)
  - Various nonlinearity, strong representability
- Disadvantage
  - Sequential process, inconvenient for parallel computing
  - High computational complexity, not suitable for high-dimensional problems with sparse features

# Outlines

Introduction

Bagging and Random Forest

Boosting and AdaBoost

Gradient Boosting Decision Tree

**XGBoost**

Conclusion and Python Examples

# Introduction

- Developed by Tianqi Chen  
(<http://homes.cs.washington.edu/~tqchen/>)
- Distributed gradient boosting : can be parallelized
- Highly efficient
- Good performance
- Out-of-Core Computing for big dataset
- Cache Optimization of data structures and algorithms

## Cost Functions

- Cost function :

$$F(\Theta_m) = \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + T(x_i; \Theta_m)) + R(\Theta_m), \text{ where}$$

$R(\Theta)$  is regularization term ( $L^0$ ,  $L^1$  or  $L^2$  penalties)

- Taylor expansion up to second order :

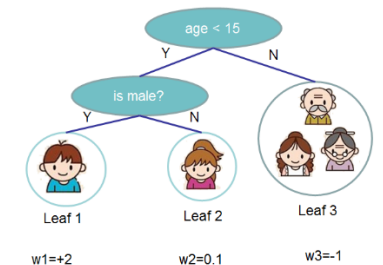
$$F(\Theta_m) \approx \sum_{i=1}^N \left[ L(y_i, f_{m-1}(x_i)) + g_i^{(m)} T(x_i; \Theta_m) + \frac{1}{2} h_{ii}^{(m)} T(x_i; \Theta_m)^2 \right] + R(\Theta_m), \text{ where}$$

$g_i^{(m)} = \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f(x_i)=f_{m-1}(x_i)}$  is the gradient of loss function, and  $h_{ii}^{(m)} = \left[ \frac{\partial^2 L(y_i, f(x_i))}{\partial f(x_i)^2} \right]_{f(x_i)=f_{m-1}(x_i)}$  is the diagonal of the Hessian of loss function (off-diagonals are zeros).



## Penalties

- Take regression trees as examples : Let  $J_m$  be the number of leaf nodes (number of rectangles in the partition),  $\gamma_{jm}$  is the approximate constant (weight  $w$ ) in the leaf node (region)  $R_{jm}$



- The complexity of tree is the sum of  $L^0$  and  $L^2$  norm of  $\{\gamma_{jm}\}$  :  $R(\Theta_m) = \frac{1}{2}\lambda \sum_{j=1}^{J_m} \gamma_{jm}^2 + \mu J_m$

$$R = \frac{1}{2}\lambda(4 + 0.01 + 1) + 3\mu$$

## Optimal solutions

- Reformulation of approximated cost function :

$$F(\Theta_m) \approx \sum_{i=1}^N L(y_i, f_{m-1}(x_i)) + \sum_{j=1}^{J_m} \left[ \left( \sum_{x_i \in R_{jm}} g_i^{(m)} \right) \gamma_{jm} + \frac{1}{2} \left( \sum_{x_i \in R_{jm}} h_{ii}^{(m)} + \lambda \right) \gamma_{jm}^2 \right] + \mu J_m =$$

$$\sum_{j=1}^{J_m} \left[ G_j^{(m)} \gamma_{jm} + \frac{1}{2} (H_j^{(m)} + \lambda) \gamma_{jm}^2 \right] + \mu J_m + \text{constant}, \text{ where}$$

$$G_j^{(m)} = \sum_{x_i \in R_{jm}} g_i^{(m)} \text{ and } H_j^{(m)} = \sum_{x_i \in R_{jm}} h_{ii}^{(m)}$$

- By differentiation w.r.t.  $\gamma_{jm}$ , we have the optimal solution :

$$\hat{\gamma}_{jm} = -\frac{G_j^{(m)}}{H_j^{(m)} + \lambda}$$

- Simplified cost function :






$$F(\Theta_m) = -\frac{1}{2} \sum_{j=1}^{J_m} \frac{(G_j^{(m)})^2}{H_j^{(m)} + \lambda} + \mu J_m + \text{constant}$$


## Structure Score

- Neglecting the constant term, we obtain the structure score :

$$SS = -\frac{1}{2} \sum_{j=1}^{J_m} \frac{(G_j^{(m)})^2}{H_j^{(m)} + \lambda} + \mu J_m$$


- It is similar to information gain : minimizing the structure score leads to the best tree

样本号	梯度数据
1 	$g_1, h_1$
2 	$g_2, h_2$
3 	$g_3, h_3$
4 	$g_4, h_4$
5 	$g_5, h_5$



## Node Splitting - Greedy Algorithm

- When splitting a node into left (L) and right (R) child nodes, we are maximizing  $Gain = \frac{1}{2} \left[ \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{(G_L + G_R)^2}{H_L + H_R + \lambda} \right] - \gamma$
- Enumerate all possible splits at  $x < a$  (e.g., age < 15) from left to right



## Greedy Algorithm for split finding

- Input : training set  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$ , loss function  $L(y, f(x))$ , the index set  $I = \{i | x_i \in R_{jm}\}$  of current node  $R_{jm}$ , feature dimension  $d$
  - Output : best split
1. Initialize  $gain = 0$ ,  $G = \sum_{i \in I} g_i$ ,  $H = \sum_{i \in I} h_{ij}$
  2. For  $k = 1$  to  $K$  :
    - 2.1  $G_L = 0$ ,  $H_L = 0$
    - 2.2 For  $j$  in sorted( $I$ , by  $x_{jk}$ ), do
      - 2.2.1  $G_L = G_L + g_j$ ,  $H_L = H_L + h_{jj}$ ,  $G_R = G - G_L$ ,  $H_R = H - H_L$
      - 2.2.2  $score = \max(score, \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{G^2}{H + \lambda})$
  3. Output split with max score

## Loss Functions

- Square loss  $L(y, f) = (y - f)^2$  :  
 $g_i^{(m)} = 2(f_i - y_i) = 2 \times \text{residue}$ ,  $h_{ii}^{(m)} = 2$
- Logistic loss  $L(y, f) = y \ln(1 + e^{-f}) + (1 - y) \ln(1 + e^f)$  :  
 $g_i^{(m)} = -y_i \left( 1 - \frac{1}{1 + e^{-f_{m-1}(x_i)}} + (1 - y_i) \frac{1}{1 + e^{-f_{m-1}(x_i)}} \right) =$   
 $Pred - Label$ ,  $h_{ii}^{(m)} = \frac{e^{-f_{m-1}(x_i)}}{(1 + e^{-f_{m-1}(x_i)})^2} = Pred \times (1 - Pred)$

# Outlines

Introduction

Bagging and Random Forest

Boosting and AdaBoost

Gradient Boosting Decision Tree

XGBoost

Conclusion and Python Examples

# Conclusions

- Ensemble methods have integrable abilities of single models, achieving better performance
- Easy to generalize to new data
- When there are strong noises, easy to overfit
- Computationally intensive



## Python Examples

- Random forest :

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators=100)
# RandomForestClassifier(bootstrap=True, class_weight=None,
    pcriterion='gini', max_depth=None, max_features='auto',
    max_leaf_nodes=None, min_impurity_split=1e-07,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    n_jobs=1, oob_score=False, random_state=None,
    verbose=0, warm_start=False)
# Feature importance in random forest
feature_imp = pd.Series(rf.feature_importances_)
rf.fit(X_train, Y_train)
Y_predict_rf = rf.predict(X_test)
oob_error = 1 - rf2.oob_score_
```

- AdaBoost :

```
from sklearn.ensemble import AdaBoostClassifier
adaboost = AdaBoostClassifier(n_estimators = 50)
adaboost.fit(X_train, Y_train)
adaboost.staged_predict(X_train)
Y_predict_ada = adaboost.predict(X_test)
```

# References

- 数据分析导论，博雅大数据学院
- 周志华，机器学习，2016
- T. Hastie, R. Tibshirani, and J. Friedman, The Elements of Statistical Learning : Data mining, Inference, and Prediction, 2nd Edition, 2009